

Reviewer guide — IPBES claim/paper verdicts

You have a file called R1_<assessment>.csv or R2_<assessment>.csv. Each row pairs one **claim** taken from an IPBES assessment with one **paper** (title and abstract). Your job is to fill in the **verdict** column.

There are two of you, working **independently**. Please do not discuss individual rows until both files are finished. We measure how often you agree, and that number is only meaningful if the judgements were made separately.

What the labels mean

Fill **verdict** with exactly one of these four words.

SUPPORTS — the abstract contains a finding that makes the claim *more likely to be true*. The paper does not have to prove the claim, cover all of it, or use the same words. It does have to report something that bears on it.

REFUTES — the abstract contains a finding that makes the claim *less likely to be true*: a contrary result, a reversed direction of effect, a failure to find what the claim asserts.

A contradiction does not have to be total. Most IPBES claims are generalisations — “X is declining”, “Y drives Z”. A paper that reports the opposite **for one region, taxon, time period or scale** makes a general claim less likely to be true, and that is **REFUTES**, not **NOT_ENOUGH_INFO**. Say in **note** how far the contradiction reaches (“only for boreal systems”, “only 1990–2005”). Use **NOT_ENOUGH_INFO** when the paper simply does not speak to the claim — not when it speaks to part of it and disagrees.

This matters because we are specifically measuring whether the system finds refutations at all. Earlier rounds produced almost none, and partial contradictions being filed as **NOT_ENOUGH_INFO** was the main reason.

NOT_ENOUGH_INFO — you have read the abstract and it does not settle the claim either way. This is the correct answer for many rows, and it is a real judgement, not a fallback. It covers:

- the paper is on a related topic but reports nothing that bears on this claim;
- the abstract is descriptive or methodological and reports no finding;
- the finding concerns a different taxon, region, driver or timescale **and nothing in it points either way** on the claim as stated.

On that last point, note the asymmetry with **REFUTES** above: a result from outside the claim’s scope that simply *does not reach* the claim is **NOT_ENOUGH_INFO**, but a result from inside part of the claim’s scope that *runs against* it is **REFUTES**. “Different scope” is not by itself a reason to withhold a verdict.

CANNOT_JUDGE — something is wrong with the *row*, not with the paper. Use it when the abstract is not in English, is truncated or garbled, is clearly not an abstract, or when the claim is too fragmentary or ambiguous to be assessed at all. Please say which in **note**.

CANNOT_JUDGE is deliberately separate from **NOT_ENOUGH_INFO**. The first means “you cannot ask me this”; the second means “I read it, and it does not answer the question”. Folding one into the other would hide a problem with how the material was assembled.

Four rules that matter more than they look

1. Judge only from the abstract in front of you. Not from the full paper, not from what you know about the topic, not from what you know about the authors. If the abstract does not say it, it is not there. This is the hardest rule to follow and the most important: the system being evaluated sees exactly this text and nothing else, so judging on extra knowledge measures the wrong thing.

2. Being on-topic is not support. A paper can be squarely about the right subject and still report nothing that bears on the specific claim. This is the single most common error, and it is the specific failure we are trying to measure — so a row that feels “obviously relevant” deserves a second look before you write **SUPPORTS**.

3. Judge the whole claim — for SUPPORTS, but not for REFUTES. Some claims bundle several assertions, or attach a mechanism to an observation. If the abstract supports one part and says nothing about the rest, that is **NOT_ENOUGH_INFO**, not **SUPPORTS** — note which part it did address.

The reverse is **not** symmetric. To support a compound claim the paper has to reach all of it; to refute one it only has to break a part, because a claim with a false component is not true as stated. One contradicted assertion out of three is **REFUTES**, with the part named in **note**.

4. The SUPPORTS / NOT_ENOUGH_INFO line. This is where reviewers and models disagree most, so one test to apply consistently: *would someone writing this claim have been able to cite this paper for it?* If the abstract reports a result you could put in a bracket after the sentence, it is **SUPPORTS**. If you would be citing it for background, framing or an adjacent fact, it is **NOT_ENOUGH_INFO**. Agreement in the same direction is not support — the paper must report a **finding**, not merely a compatible view.

Filling in the file

column	what to put
verdict	one of the four words above — required
note	free text; always worth a line when you hesitated, and required for CANNOT_JUDGE

column	what to put
reviewer	your initials — same value on every row
date	the date you judged the row, YYYY-MM-DD

Edit **verdict**, **note**, **reviewer** and **date** only. Please do not reorder, add or delete rows, and do not change **id** — the two files are matched on it. The row order differs between the two of you on purpose.

The file is plain CSV. If you open it in Excel, save it back as CSV (UTF-8), not as **.xlsx**.

Pace, and when to stop

Expect around 1.5 minutes per row, so roughly five hours for 200 rows. Please spread it over several sittings — agreement drops measurably when reviewers tire, and the comparison is between the two of you.

If you find yourself unsure for more than a minute or so, write the verdict you lean towards and say so in **note**. A recorded hesitation is more useful than a confident guess, and disagreements between you are something we want to see, not something to avoid.

What happens next

The two files are compared row by row. Rows where you agree become the gold standard that every model is scored against. Rows where you disagree are reported as the *human ceiling* — the limit of how well any automated system could be expected to do — and are not used to score models.